

EXPERIMENT 10

AIM: Implementing Resource Quota in Kubernetes

Objective:

In Kubernetes, Resource Quotas are used to control the resource consumption of namespaces. They help in managing and enforcing limits on the usage of resources like CPU, memory, and the number of objects (e.g., Pods, Services) within a namespace. This exercise will guide you through creating and managing Resource Quotas to limit the resources used by applications in a specific namespace.

Step 1: Understand Resource Quotas

Resource Quotas allow you to:

- Limit the amount of CPU and memory a namespace can use.
- Control the number of certain types of resources (e.g., Pods, Services, PersistentVolumeClaims) in a namespace.
- Prevent a namespace from consuming more resources than allocated, ensuring fair usage across multiple teams or applications.

Step 2: Create a Namespace

First, create a namespace where you will apply the Resource Quota. This helps in isolating and controlling resource usage within that specific namespace.

Create a YAML file named quota-namespace.yaml with the following content:

The screenshot shows a terminal window titled "aryanbansal — nano quota-namespace.yaml — 155x52". The file content is as follows:

```
apiVersion: v1
kind: Namespace
metadata:
  name: quota-example  # The name of the namespace.
```

The terminal window has a dark background and light-colored text. At the bottom, there is a menu bar with various keyboard shortcuts for file operations like Get Help, WriteOut, Read File, Prev Pg, Next Pg, Cut Text, Uncut Text, Cur Pos, and To Spell.

Apply the YAML to create the namespace:

```
kubectl apply -f quota-namespace.yaml
```

```
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl apply -f quota-namespace.yaml
namespace/quota-example created
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

Verify that the namespace is created:

```
kubectl get namespaces
```

```
[(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get namespaces
NAME        STATUS  AGE
default     Active  11d
kube-node-lease  Active  11d
kube-public   Active  11d
kube-system   Active  11d
quota-example Active  35s
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % ]
```

You should see quota-example listed in the output.

Step 3: Define a Resource Quota

Next, create a Resource Quota YAML file named resource-quota.yaml with the following content:



The screenshot shows a terminal window with the title bar "aryanbansal — nano resource-quota.yaml — 155x52". The file content is as follows:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: example-quota  # The name of the Resource Quota.
  namespace: quota-example # The namespace to which the Resource Quota will apply.
spec:
  hard:
    requests.cpu: "2"      # The hard limits imposed by this Resource Quota.
    requests.memory: "4Gi" # The total CPU resource requests allowed in the namespace (2 cores).
    limits.cpu: "4"         # The total CPU resource limits allowed in the namespace (4 cores).
    limits.memory: "8Gi" # The total memory resource limits allowed in the namespace (8 GiB).
    pods: "10"            # The total number of Pods allowed in the namespace.
    persistentvolumeclaims: "5" # The total number of PersistentVolumeClaims allowed in the namespace.
    configmaps: "10"        # The total number of ConfigMaps allowed in the namespace.
    services: "5"          # The total number of Services allowed in the namespace.
```

The bottom of the terminal window shows nano editor key bindings:

$\wedge G$ Get Help	$\wedge O$ WriteOut	$\wedge R$ Read File	$\wedge Y$ Prev Pg	$\wedge K$ Cut Text	$\wedge Q$ Cur Pos
$\wedge X$ Exit	$\wedge J$ Justify	$\wedge W$ Where is	$\wedge V$ Next Pg	$\wedge U$ UnCut Text	$\wedge T$ To Spell

Step 4: Apply the Resource Quota

Apply the Resource Quota YAML to the namespace:

```
kubectl apply -f resource-quota.yaml
```

```
[(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl apply -f resource-quota.yaml
resourcequota/example-quota created
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % ]
```

Verify that the Resource Quota is applied:
kubectl get resourcequota -n quota-example

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get resourcequota -n quota-example
NAME          AGE   REQUEST           LIMIT
example-quota 2m6s  configmaps: 1/10, persistentvolumeclaims: 0/5, pods: 0/10, requests.cpu: 0/2, requests.memory: 0/4Gi, services: 0/5  limits.cpu: 0/4, limits.memory: 0/8Gi
Name:          example-quota
Namespace:    quota-example
Resource       Used   Hard
configmaps    1     10
limits.cpu    0     4
limits.memory 0     8Gi
persistentvolumeclaims 0     5
pods          0     10
requests.cpu  0     2
requests.memory 0     4Gi
services      0     5
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

To describe the Pods and see their resource allocations:
kubectl describe pods -l app=nginx -n quota-example

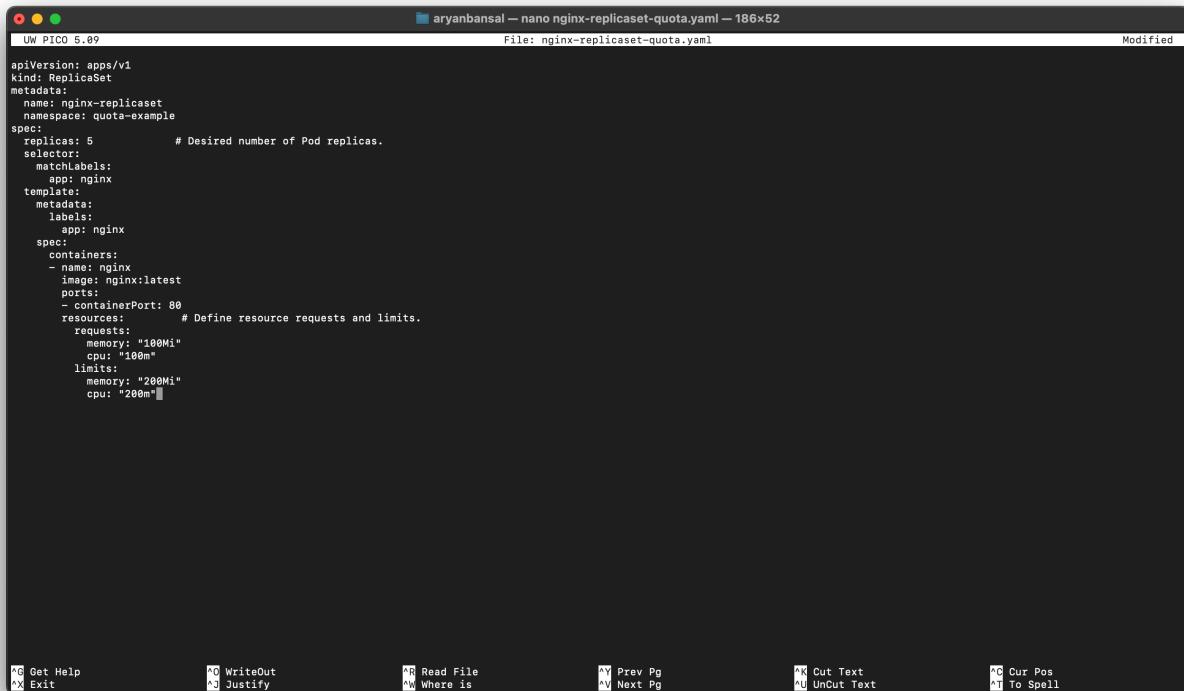
```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl describe pods -l app=nginx -n quota-example
No resources found in quota-example namespace.
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

Step 5: Test the Resource Quota

Let's create some resources in the quota-example namespace to see how the Resource Quota affects them.

Deploy a ReplicaSet with Resource Requests and Limits

Create a YAML file named nginx-replicaset-quota.yaml with the following content:



```
aryanbansal — nano nginx-replicaset-quota.yaml — 186x52
File: nginx-replicaset-quota.yaml
Modified

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
  namespace: quota-example
spec:
  replicas: 5          # Desired number of Pod replicas.
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
          resources:
            requests:
              memory: "100Mi"
              cpu: "10m"
            limits:
              memory: "200Mi"
              cpu: "20m"
```

Explanation:

This ReplicaSet requests a total of 500m CPU and 500Mi memory across 5 replicas.

It also limits each replica to use a maximum of 200m CPU and 200Mi memory.

Apply this YAML to create the ReplicaSet:

```
kubectl apply -f nginx-replicaset-quota.yaml
```

```
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl apply -f nginx-replicaset-quota.yaml
replicaset.apps/nginx-replicaset created
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

Check the status of the Pods and ensure they are created within the constraints of the Resource Quota:

```
kubectl get pods -n quota-example
```

```
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get pods -n quota-example
NAME          READY   STATUS    RESTARTS   AGE
nginx-replicaset-2fcckj  1/1    Running   0          32s
nginx-replicaset-t5zwx  1/1    Running   0          32s
nginx-replicaset-t6ww9  1/1    Running   0          32s
nginx-replicaset-v8dwk  1/1    Running   0          32s
nginx-replicaset-vgrzz  1/1    Running   0          32s
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

To describe the Pods and see their resource allocations:

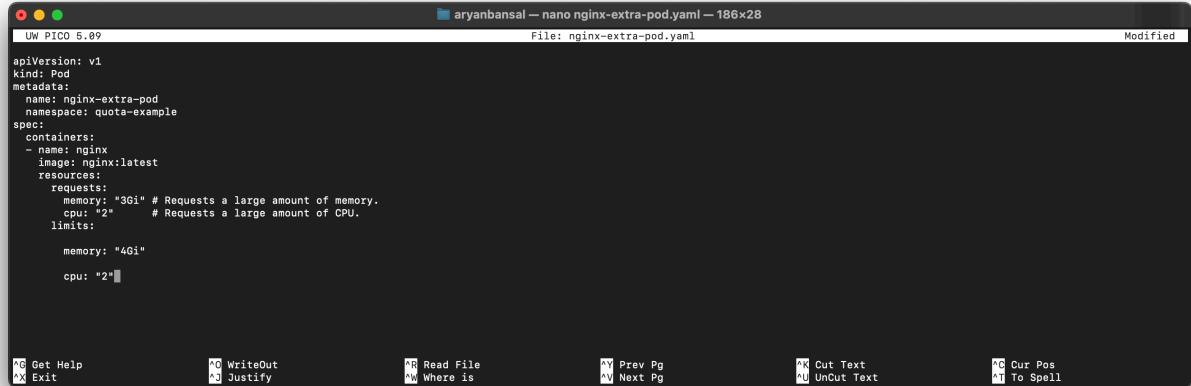
```
kubectl describe pods -l app=nginx -n quota-example
```

```
aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl apply -f resource-quota.yaml
resourcequota/example-quota created
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get resourcequotas -n quota-example
NAME           AGE   REQUEST                                LIMIT
example-quota  41s   configmaps: 1/10, persistentvolumeclaims: 0/5, pods: 0/10, requests.cpu: 0/2, requests.memory: 0/4Gi, services: 0/5   limits.cpu: 0/4, limits.memory: 0/8Gi
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl describe resourcequotas example-quota -n quota-example
No resources found in quota-example namespace.
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get resourcequotas -n quota-example
NAME           AGE   REQUEST                                LIMIT
example-quota  2m6s  configmaps: 1/10, persistentvolumeclaims: 0/5, pods: 0/10, requests.cpu: 0/2, requests.memory: 0/4Gi, services: 0/5   limits.cpu: 0/4, limits.memory: 0/8Gi
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl describe resourcequotas example-quota -n quota-example
Name:           example-quota
Namespace:      quota-example
Resource        Used Hard
-----
configmaps     1    10
limits.cpu     0    4
limits.memory  0    8Gi
persistentvolumeclaims 0    5
pods           0    10
requests.cpu   0    2
requests.memory 0    4Gi
services        0    5
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl describe pods -l app=nginx -n quota-example
No resources found in quota-example namespace.
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % touch nginx-replicaset-quota.yaml
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % nano nginx-replicaset-quota.yaml
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl apply -f nginx-replicaset-quota.yaml
replicaset.apps/nginx-replicaset created
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get pods -n quota-example
NAME          READY   STATUS    RESTARTS   AGE
nginx-replicaset-2fcckj  1/1    Running   0          32s
nginx-replicaset-t5zwx  1/1    Running   0          32s
nginx-replicaset-t6ww9  1/1    Running   0          32s
nginx-replicaset-v8dwk  1/1    Running   0          32s
nginx-replicaset-vgrzz  1/1    Running   0          32s
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl describe pods -l app=nginx -n quota-example
Name:           nginx-replicaset-2fcckj
Namespace:      quota-example
Priority:       0
Service Account: default
Node:          minikube/192.168.49.2
Start Time:    Thu, 21 Nov 2024 19:30:44 +0530
Labels:         app=nginx
Annotations:   <none>
Status:        Running
IP:            10.244.0.24
IPs:
  IP:          10.244.0.24
Controlled By: ReplicaSet/nginx-replicaset
Containers:
  nginx:
    Container ID:  docker://e5dc13f3fee5ccb83a4784b1f129355277b1f6bfe421c58dad694429f53ac09c
    Image:         nginx:latest
    Image ID:     docker-pullable://nginx@sha256:bc5eacf581eda30808b4b1f07ebba230de2f27d47767129a6a905c84f470
    Port:         80/TCP
    Host Port:   0/TCP
```

Attempt to Exceed the Resource Quota

Try creating additional resources to see if they are rejected when exceeding the quota. For example, create more Pods or increase the CPU/memory requests to exceed the quota limits.

Create a YAML file named nginx-extra-pod.yaml with the following content:



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-extra-pod
  namespace: quota-example
spec:
  containers:
    - name: nginx
      image: nginx:latest
      resources:
        requests:
          memory: "3Gi" # Requests a large amount of memory.
          cpu: "2" # Requests a large amount of CPU.
        limits:
          memory: "4Gi"
          cpu: "2"
```

Apply this YAML to create the Pod:

```
kubectl apply -f nginx-extra-pod.yaml
```

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl apply -f nginx-extra-pod.yaml
pod/nginx-extra-pod created
```

This should fail due to exceeding the Resource Quota. Check the events to see the failure reason:
kubectl get events -n quota-example

Look for error messages indicating that the Pod creation was denied due to resource constraints..

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get events -n quota-example
LAST SEEN   TYPE    REASON   OBJECT            MESSAGE
2m15s     Normal  Scheduled  pod/nginx-extra-pod  Successfully assigned quota-example/nginx-extra-pod to minikube
2m14s     Normal  Pulling   pod/nginx-extra-pod  Pulling image "nginx:latest"
2m8s      Normal  Pulled    pod/nginx-extra-pod  Successfully pulled image "nginx:latest" in 6.021s (6.021s including waiting). Image size: 196880043 bytes.
2m8s     Normal  Created   pod/nginx-extra-pod  Created container nginx
2m8s     Normal  Started   pod/nginx-extra-pod  Started container nginx
2m5s      Normal  Killing   pod/nginx-extra-pod  Stopping container nginx
44s      Normal  Scheduled  pod/nginx-extra-pod  Successfully assigned quota-example/nginx-extra-pod to minikube
43s      Normal  Pulling   pod/nginx-extra-pod  Pulling image "nginx:latest"
38s      Normal  Pulled    pod/nginx-extra-pod  Successfully pulled image "nginx:latest" in 4.846s (4.846s including waiting). Image size: 196880043 bytes.
38s      Normal  Created   pod/nginx-extra-pod  Created container nginx
38s      Normal  Started   pod/nginx-extra-pod  Started container nginx
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

Step 6: Clean Up Resources

To delete the resources you created:

```
kubectl delete -f nginx-replicaset-quota.yaml
```

```
kubectl delete -f nginx-extra-pod.yaml
```

```
kubectl delete -f resource-quota.yaml
```

```
kubectl delete namespace quota-example
```

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl delete -f nginx-replicaset-quota.yaml
replicaset.apps "nginx-replicaset" deleted
```

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl delete -f nginx-extra-pod.yaml
Error from server (NotFound): error when deleting "nginx-extra-pod.yaml": pods "nginx-extra-pod" not found
```

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl delete -f resource-quota.yaml
```

```
[kubect] delete namespace quota-example
```

```
resourcequota "example-quota" deleted
```

```
namespace "quota-example" deleted
```

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl delete -f nginx-extra-pod.yaml
```

```
pod "nginx-extra-pod" deleted
```

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

